

N I S O L A I

AI Transformation, Delivered.

ENTERPRISE AI TRANSFORMATION ENGAGEMENT

Board-Ready AI Strategy, Capability Assessment & Roadmap

PREPARED FOR

Novatech Systems Ltd.

INDUSTRY SECTOR

**Technology, SaaS & IT
Services**

PREPARED BY

Nisol AI Advisory

DATE OF ISSUE

August 31, 2026

CLASSIFICATION

Commercial-in-Confidence

DOCUMENT ID

**Novatech-Systems-Ltd--AI-
Transformation-2026-01**

TABLE OF CONTENTS

1. Executive Summary & Value Proposition	3
1.1 Engagement Overview 1.2 Key Drivers 1.3 Expected Business Impact 1.4 Executive Sidebar	
2. AI Readiness, Risk & Data Assessment	7
2.1 Maturity Score vs Sector 2.2 Capability Radar 2.3 Department Heatmap 2.4 Vulnerabilities 2.5 Risk Register 2.6 Data Readiness	
3. AI Opportunity Matrix & Use Cases Catalog	12
3.1 Matrix Visualization 3.2 Prioritized Use Cases Catalog 3.3 Quick Wins vs Strategic Bets	
4. Transformation Roadmap & Change Management (OCM)	16
4.1 Phased Implementation Timeline 4.2 Change Management & RACI Matrix 4.3 Training & Adoption Plan	
5. ROI Analysis, Sensitivity & Financial Projection	19
5.1 5-Year Benefit vs Investment 5.2 Department Breakdown 5.3 5-Year Projection Table 5.4 3-Scenario Sensitivity Stress Test	
6. Solution Blueprints	23
Technical Architecture & Engineering Specifications for High-Priority Initiatives	
7. Implementation Investment & Commercial	26
7.1 Prioritized Initiatives Table 7.2 Engagement Tiers 7.3 Next Steps	
8. Terms & Conditions & Executive Acceptance	30
8.1 Commercial Terms 8.2 Executive Authorization Sign-off	

1. Executive Summary & Value Proposition

1.1 Engagement Overview

Novatech Systems Ltd. is an established leader in the **Technology, SaaS & IT Services** sector preparing for accelerated operating scale. Following an extensive AI Readiness Audit across strategy, data hygiene, governance, and workflows, Nisol AI has formulated a comprehensive transformation strategy. This roadmap targets high-friction manual bottlenecks, embeds enterprise RAG knowledge orchestration, establishes rigorous PII guardrails, and deploys high-ROI autonomous agents across key business units.

1.2 Expected Business Impact

OVERALL AI
READINESS

34%

Industry Avg: 62%

HIGH PRIORITY
QUICK WINS

10

10 Strategic Bets

EST. ANNUAL
SAVINGS

₹4.8

Crore

Across 7
Departments

EST. PAYBACK
PERIOD

8

Months

Avg. ROI: 340%

1.3 EXECUTIVE ENGAGEMENT SUMMARY — AT A GLANCE

Client Organization: Novatech Systems Ltd.

Industry Sector: Technology, SaaS & IT Services

Overall Maturity Score: 1.7 / 5.0 (34%)

Sector Median Benchmark: 3.6 / 5.0 (72%)

Benchmark Deficit: 38 pts behind sector median

Sector Top Quartile: 4.5 / 5.0 (91%)

Initial Program Investment: ₹4.26 Crore

Estimated Annual Savings: ₹8.80 Crore

5-Year Cumulative Net Benefit: ₹36.49 Crore

Expected Payback Period: 5.8 Months

Estimated 5-Year Program ROI: +634%

* Benchmark Citation: Nisol AI Enterprise Benchmark Index (Q3 2026), n=140+ tech enterprises, calibrated against Gartner & Stanford AI Index maturity frameworks.


1.4 Strategic Imperative: Cost of Inaction & Valuation Multiple Impact

The Cost of Delay: Operating in the **Technology, SaaS & IT Services** sector without enterprise AI automation introduces an estimated **₹8.80 Crore** annual drag in avoidable manual labor and ticket cycle friction. Top-quartile competitors reinvest automation savings into product development cycles that outpace peers 3:1. Furthermore, public markets and private equity sponsors reward AI-leveraged operating models with an average **15%–25% valuation multiple expansion** due to higher margin defensibility.

CONFIDENTIAL

2. AI Readiness, Risk & Data Assessment

2.1 Maturity Score vs. Technology, SaaS & IT Services Benchmark




Root Cause Analysis of the 38-Point Benchmark Gap

Novatech Systems Ltd.'s overall maturity score of **1.7 / 5.0 (34%)** places the enterprise 38 points behind the Technology, SaaS & IT Services median (3.6 / 5.0). This deficit is primarily driven by ad-hoc departmental experimentation lacking centralized API security gateways (-7.2 pts), unstructured document silos in finance and operations (-6.4 pts), and manual regression testing (-6.4 pts). Wave 1 and Wave 2 implementations target these root causes directly to close 14 of the 20 gap points within 180 days.

2.2 Capability Maturity Radar (8 Strategic Dimensions)

Capability Maturity Radar (8 Dimensions)

Assessment Score



2.3 Department AI Capability Heatmap

Department AI Capability Heatmap

4.0-5.0 High 3.0-3.9 Mid 2.2-2.9 Low

	Data Quality	Automation	Tooling	Governance
Sales & Business Development	1.9 / 5	1.5 / 5	1.9 / 5	1.5 / 5
Marketing & Growth	1.9 / 5	1.5 / 5	1.9 / 5	1.5 / 5
Finance & Accounting	1.9 / 5	1.5 / 5	1.9 / 5	1.5 / 5
Customer Support & CX	1.9 / 5	1.5 / 5	1.9 / 5	1.5 / 5
Software Engineering	1.9 / 5	1.5 / 5	1.9 / 5	1.5 / 5
IT Operations & Cloud	1.9 / 5	1.5 / 5	1.9 / 5	1.5 / 5

2.4 Departmental Action Briefs & Governance Playbook


Functional relevance briefs establishing operational friction points, quantifiable target outcomes, Hub-and-Spoke executive ownership, and low-friction SME engagement models across audited enterprise units.

Department & Sponsor	Current Operational Friction	Target AI Outcome	Governance & SME Ask
Sales & Business Development Chief Commercial Officer / VP Sales Developing Baseline	Reps spend 20+ hrs/wk drafting complex RFP bids and researching prospect firmographic intent manually.	RAG Proposal Generator cuts bid turnaround from 3 days to 25 mins; Predictive lead scoring boosts conversion by +22%.	Model: Sales VP controls pricing logic and margin guardrails; Nisol AI delivers the secure RAG proposal engine. SME Ask: 1 Senior Sales Rep for 3 hrs/wk during Week 8 to calibrate winning proposal examples and rate cards.
Marketing & Growth Chief Marketing Officer (CMO) Developing Baseline	80+ hrs/wk spent manually drafting and formatting multi-channel A/B ad variations across regional segments.	Generative Content Engine synthesizes 50 localized variants in seconds; +15% conversion lift without agency spend.	Model: CMO retains 100% creative control; AI functions as a high-velocity drafting accelerator. SME Ask: 1 Content Lead for 2 hrs/wk to calibrate brand tone-of-voice and style guidelines.
Finance & Accounting Chief Financial Officer / Corporate Controller Developing Baseline	15,000 invoices/month with 5 manual touches each; 10-day month-end financial close cycle across ledgers.	Multi-modal OCR cuts touches from 5 to 1; Real-time reconciliation compresses close cycle from 10 days to 4 days.	Model: Finance Controller retains full approval thresholds and exception signing authority. SME Ask: 1 Accounting Lead for 4 hrs/wk in Weeks 8–10 to review sample invoice extractions in staging sandbox.
Customer Support & CX VP Customer Experience / Support Developing Baseline	High agent burnout; 40% of tier-1 tickets are repetitive password, billing, and status inquiries.	In-line Agent Co-Pilot auto-drafts responses from historical ticket archives, reducing AHT by 40%.	Model: Support VP reviews response templates; human-in-the-loop agent approves before sending. SME Ask: 2 Senior Agents for 3 hrs in sandbox to validate draft suggestion accuracy.
Software Engineering VP Engineering / CTO Developing Baseline	Senior engineers spend 15 hrs/wk writing boilerplate regression suites and conducting manual code reviews.	AI agent auto-generates Playwright/Cypress end-to-end tests, recovering 10 dev hours/week per senior engineer.	Model: Engineering VP retains full git merge controls; AI runs strictly as a CI/CD checking pipeline. SME Ask: 1 QA Lead for 4 hrs/wk to review initial 20 auto-generated Playwright tests.

Department & Sponsor	Current Operational Friction	Target AI Outcome	Governance & SME Ask
IT Operations & Cloud Chief Information Officer / VP IT Operations Developing Baseline	6-hour MTTR on server alerts; manual VPN and SaaS ticket provisioning backlogs.	Autonomous IT Helpdesk Bot resolves 55% of L1 access requests in <60 seconds, cutting ticket backlog by 45%.	Model: CIO sets access boundaries and SSO roles; bot operates under strict least-privilege tokens. SME Ask: 1 Sysadmin for 2 hrs/wk to test Jira/Slack provisioning action scripts.
InfoSec & Data Privacy Chief Information Security Officer (CISO) / DPO Developing Baseline	Zero visibility into employee shadow AI usage; high risk of proprietary IP and customer PII exfiltration.	Zero-retention enterprise proxy with inline PII masking, guaranteeing 100% audit logging.	Model: CISO controls DLP regex patterns and compliance inspection logs. SME Ask: 1 SecOps Analyst for 2 hrs to review tokenization proxy architecture and audit logging.
Legal & Compliance General Counsel / Chief Compliance Officer Developing Baseline	200+ vendor NDAs and master service agreements reviewed monthly with 5-7 business day backlog.	RAG contract review highlights non-standard clauses and flags deviations in <5 minutes, cutting turnaround by 70%.	Model: General Counsel establishes fallback clause taxonomy; AI never signs or executes contracts. SME Ask: 1 In-house Counsel for 3 hrs in Week 8 to calibrate acceptable standard clauses.

2.5 Enterprise AI Risk & Regulatory Register

Prioritized risk matrix mapping technical, regulatory (SOC 2 Type II, ISO 27001, DPDP Act 2023), and operational vulnerabilities with concrete mitigation strategies.



Risk ID	Category & Description	Severity	Mitigation Strategy	Accountability
RSK-01	Data Privacy & Security PII or sensitive client intellectual property inadvertently sent to third-party public LLM endpoints without token redaction.	20/25 (Critical)	Deploy an enterprise API Gateway proxy with regex + Presidio automated PII redaction layer prior to external payload dispatch.	Chief Information Security Officer (CISO)
RSK-02	Model Risk & Bias LLM hallucinations or biased outputs generated in customer-facing conversational interfaces or financial reports.	16/25 (High)	Enforce strict temperature clamping (≤ 0.2), grounded RAG with vector citation validation, and Human-in-the-Loop (HITL) sign-off on critical workflows.	Head of AI Engineering & QA
RSK-03	Regulatory & Compliance Non-compliance with data localization mandates and emerging AI auditing standards.	12/25 (High)	Host vector databases and self-managed open-weights LLMs within in-region private cloud VPCs (e.g. AWS Mumbai / Azure India).	Chief Legal Officer / Compliance Head

Risk ID	Category & Description	Severity	Mitigation Strategy	Accountability
RSK-04	Operational & Adoption	12/25 (Medium)	Roll out sanctioned enterprise AI workspace with single-sign-on (SSO), comprehensive staff upskilling, and active department champions.	VP Human Resources & Change Management Lead
	Staff resistance or shadow AI adoption using unauthorized consumer AI tools without corporate audit trails.			
RSK-05	Vendor & Infrastructure	9/25 (Medium)	Architect model-agnostic abstraction layer (LiteLLM / OpenRouter architecture) enabling seamless multi-model fallback.	Principal Enterprise Architect
	Excessive dependency on a single proprietary LLM provider risking pricing spikes, API latency degradation, or deprecation.			

CONFIDENTIAL

2.6 Data Strategy & Readiness Assessment

Data is the foundational determinant of AI ROI. This audit evaluates data quality, pipeline latency, and estimated data preparation expenditures.

Data Quality Dimensions

Completeness	57% (Needs Attention)
Accuracy	61% (Needs Attention)
Timeliness	53% (Healthy)
Consistency	51% (Needs Attention)
Accessibility	55% (Critical Blocker)

Estimated Data Preparation Budget


₹28 - ₹38 Lakhs

Accounting for ~42% of Phase 1 implementation allocation for vector lakehouse ETL & PII scrubbing.

CONFIDENTIAL

3. AI Opportunity Matrix & Use Cases Catalog

3.1 Opportunity Matrix (Value vs. Implementation Effort)



3.2 Prioritized Top AI Use Cases Catalog (Phased Execution Horizons)

Prioritized portfolio mapping 20 cross-functional enterprise initiatives across 6 execution horizons, balancing immediate time-to-value with long-term competitive moat creation under the Single-Pod delivery benchmark.

#	Initiative & Execution Horizon	Category & Pod	Business Problem & Proposed Solution	Duration, Savings & ROI
1	AI-Driven Automated QA Test Case Generation & Execution Horizon 1: Foundation (M 0-3) Software Engineering & QA	Quick Win 1 Dedicated Pod (AI Lead + 2 Devs + 0.5 QA) Complexity: Low	Pain Point: Manual regression testing of Novatech's enterprise software platforms delays sprint releases by 48-72 hours, creating developer bottlenecks and regression bugs. Solution: Autonomous agent analyzing PR code diffs to auto-generate end-to-end Playwright/Cypress tests pre-merge.	+280% ROI ₹32 Lakhs/year (1,400 engineering hours saved) Duration: 5 Weeks
2	AI Proposal & Technical RFP Response Generator	Quick Win	Pain Point: Drafting technical RFP bids for Novatech's software services takes 20+ hours per	+340% ROI

#	Initiative & Execution Horizon	Category & Pod	Business Problem & Proposed Solution	Duration, Savings & ROI
	Horizon 1: Foundation (M 0-3) Sales & Business Development	1 Dedicated Pod Complexity: Low	submission, utilizing scattered knowledge across SharePoint and legacy CRMs. Solution: RAG-powered bid responder indexing past winning proposals, security audits, and rate cards to generate 80% drafts in minutes.	₹38 Lakhs/year (50% faster proposal turnaround) Duration: 6 Weeks
3	Autonomous PMO Operational Status & Flight-Deck Reporter Horizon 1: Foundation (M 0-3) Corporate Strategy & PMO	Quick Win 1 Dedicated Pod Complexity: Low	Pain Point: Novatech project managers spend 6-8 hours weekly manually pulling project status across Jira, Slack, and SharePoint to assemble executive slide decks. Solution: Autonomous PMO agent synthesizing weekly sprint velocities, milestone blockers, and burn rates into formatted executive digests.	+220% ROI ₹18 Lakhs/year (600 PM hours recovered) Duration: 4 Weeks
4	Intelligent Multi-Modal Invoice & Vendor Document OCR Processing Horizon 2: Dev Acceleration (M 3-6) Finance & Accounting	Quick Win 1 Dedicated Pod Complexity: Medium	Pain Point: Processing 15,000 monthly supplier invoices at Novatech involves manual data extraction and matching against ERP ledgers, causing late payments and compliance delays. Solution: Vision LLM extracting 15 line items, matching PO numbers against ERP ledgers, and routing approved batches for payment.	+310% ROI ₹42 Lakhs/year in labor and early-payment discounts Duration: 10 Weeks
5	Enterprise Knowledge Graph & Cross-Repo Developer Search Engine Horizon 2: Dev Acceleration (M 3-6) Engineering & IT	Strategic Bet 1 Dedicated Pod Complexity: High	Pain Point: Novatech engineers and product managers lose 4+ hours weekly searching across siloed SharePoint drives, Confluence wikis, and legacy code repos for architectural patterns. Solution: Unified hybrid vector + knowledge graph search engine delivering instant verified code snippets and system diagrams.	+360% ROI ₹65 Lakhs/year in recovered developer velocity Duration: 20 Weeks
6	Natural Language Text-to-SQL Enterprise BI Query Agent Horizon 3: Operational Scale (M 6-9) Data & Analytics	Strategic Bet 1 Dedicated Pod + 1 Data Eng Complexity: Medium	Pain Point: Novatech business managers submit 50+ ad-hoc SQL query requests weekly, creating a 2-week backlog for the data team due to siloed SQL Data Warehouses. Solution: Semantic data agent allowing non-technical managers to query the data warehouse via natural English with SQL validation.	+310% ROI ₹34 Lakhs/year (65% reduction in ad-hoc query backlog) Duration: 14 Weeks
7	Enterprise Zero-Trust PII Tokenization Reverse	Strategic Bet	Pain Point: Novatech lacks centralized AI governance, creating high compliance risks of accidental	+310% ROI

#	Initiative & Execution Horizon	Category & Pod	Business Problem & Proposed Solution	Duration, Savings & ROI
	Proxy Horizon 3: Operational Scale (M 6-9) InfoSec & Privacy	1 Dedicated Pod Complexity: Medium	PII/source code exfiltration as employees adopt public LLM tools. Solution: Network-level reverse proxy intercepting LLM prompts to tokenize sensitive entities locally before cloud egress.	₹50 Lakhs/year in compliance risk avoidance Duration: 12 Weeks
8	Automated IT Helpdesk L1 Ticket Resolution & Provisioning Agent Horizon 3: Operational Scale (M 6-9) IT Operations & Cloud	Quick Win 1 Dedicated Pod Complexity: Medium	Pain Point: Employees wait hours for routine VPN resets, SaaS seat provisioning, and hardware access approvals. Solution: Autonomous agent resolving 55% of L1 IT requests via authenticated Slack commands in sub-60 seconds.	+275% ROI ₹24 Lakhs/year in IT support labor Duration: 7 Weeks
9	Omnichannel Customer Support Agent Co-Pilot & Ticket Assist Horizon 4: CX & Pipeline (M 9-12) Customer Service & CX	Quick Win 1 Dedicated Pod + Support Lead Complexity: Medium	Pain Point: High tier-1 ticket volume causes long resolution times and frontline agent burnout during peak volume spikes. Solution: In-line Agent Co-Pilot auto-drafting responses from past ticket resolutions for human one-click review and send.	+295% ROI ₹36 Lakhs/year (40% reduction in AHT) Duration: 12 Weeks
10	Predictive Lead Scoring & Multi-Touch Attribution Engine Horizon 4: CX & Pipeline (M 9-12) Sales & Marketing	Quick Win 1 Dedicated Pod Complexity: Medium	Pain Point: Sales reps waste time prospecting unqualified leads while high-intent inbound inquiries go cold. Solution: Machine learning intent tracker scoring leads based on email velocity, website behavior, and firmographic signals.	+350% ROI ₹45 Lakhs/year in pipeline conversion uplift (+22%) Duration: 10 Weeks
11	Automated Marketing Content Engine & Campaign Localization Horizon 4: CX & Pipeline (M 9-12) Marketing & Growth	Quick Win 1 Dedicated Pod Complexity: Low	Pain Point: Creating localized ad copy and content variants across multiple regional languages and personas requires 80+ hours weekly. Solution: Brand-aligned generative content engine synthesizing 50+ validated ad variants with automated brand compliance checks.	+230% ROI ₹20 Lakhs/year in external agency fees Duration: 6 Weeks
12	Enterprise Knowledge Graph & Cross-Repo Developer Search Engine Horizon 5: Strategic Bets (M 12-18) Engineering & IT	Strategic Bet 2 Parallel Pods (Lead Architect + 4 Devs + 2 Data Engs) Complexity: High	Pain Point: Engineers lose 4 hours weekly searching siloed repos, Confluence wikis, and Slack threads for architectural patterns. Solution: Unified hybrid vector + knowledge graph search engine delivering instant verified code snippets and system diagrams.	+360% ROI ₹65 Lakhs/year in recovered developer velocity Duration: 20 Weeks

#	Initiative & Execution Horizon	Category & Pod	Business Problem & Proposed Solution	Duration, Savings & ROI
13	Dynamic Price Optimization & Margin Analytics Engine Horizon 5: Strategic Bets (M 12-18) Finance & Strategy	Strategic Bet 2 Parallel Pods Complexity: High	Pain Point: Pricing models rely on static quarterly spreadsheets, failing to capture raw material fluctuations and competitor changes. Solution: Real-time algorithmic pricing engine simulating demand elasticity and margin impact across product lines.	+380% ROI ₹55 Lakhs/year in captured margin leakage Duration: 18 Weeks
14	Natural Language Text-to-SQL Enterprise BI Query Agent Horizon 5: Strategic Bets (M 12-18) Data & Analytics	Strategic Bet 1 Dedicated Pod + 1 Data Eng Complexity: Medium	Pain Point: Business managers submit 50+ ad-hoc SQL query requests weekly, creating a 2-week backlog for data analysts. Solution: Semantic data agent allowing non-technical managers to query the data warehouse via natural English with SQL validation.	+310% ROI ₹34 Lakhs/year (65% reduction in ad-hoc query backlog) Duration: 14 Weeks
15	Predictive Supply Chain Demand & Inventory Lead-Time Optimizer Horizon 5: Strategic Bets (M 12-18) Operations & Logistics	Strategic Bet 2 Parallel Pods Complexity: High	Pain Point: Buffer stock over-allocations and sudden component stockouts tie up working capital in static warehouses. Solution: Time-series forecasting model balancing real-time sales pipeline velocity against international supplier shipping times.	+410% ROI ₹72 Lakhs/year in reduced inventory carrying costs Duration: 22 Weeks
16	Autonomous Infrastructure Monitoring & Self-Healing SRE Ops Horizon 6: Autonomous Ops (M 18-24) IT Operations & Cloud	Strategic Bet 2 Dedicated SRE Pods + Client SRE Leads Complexity: High	Pain Point: Production incidents require manual triage across logs, Datadog metrics, and runbooks, risking client SLA penalties. Solution: Autonomous agent correlating telemetry across server clusters, diagnosing root causes, and executing verified remediation runbooks.	+420% ROI ₹85 Lakhs/year in avoided downtime and on-call toil Duration: 26 Weeks

4. Transformation Roadmap & Change Management (OCM)

4.1 Transformation Roadmap (Single-Pod Execution Waves)

Governed by Nisol AI's **Single-Pod Capacity Benchmark**: exactly 1 dedicated pod executes 1 primary strategic initiative per 12-to-14 week wave. Concurrent initiative delivery requires provisioning independent parallel pods.

Wave 1 (M 0-3): Foundation & Anchor Quick Win

Deploy single anchor high-ROI automation (1 Pod) and establish enterprise AI policy, zero-retention VPC gateway, and baseline data hygiene.

Core Initiatives: AI-Driven Automated QA Test Case Generation & Execution; Enterprise AI Security Proxy & Acceptable Use Policy Setup

Investment: ₹18 - ₹25 Lakhs (Professional Services) |
Owner: Head of AI Engineering & Primary Dept Sponsor

Wave 2 (M 3-6): Dev Acceleration & Department Scaling

Stabilize Wave 1 in production and execute second prioritized initiative (1 Pod). Concurrency requires independent parallel pod.

Core Initiatives: Intelligent Multi-Modal Invoice & Vendor Document OCR Processing; Wave 1 Canary Production Handover

Investment: ₹22 - ₹32 Lakhs (Professional Services) |
Owner: VP of Product / Business Unit Leads

Wave 3 (M 6-9): Operational Scale & Financial Automation

Scale AI workflow into business operations (Finance/Operations) with live ERP/CRM read-only connectors.

Core Initiatives: Predictive Lead Scoring & Multi-Touch Attribution Engine; Data Lakehouse ETL Ingestion Pipeline

Investment: ₹28 - ₹38 Lakhs (Professional Services) |
Owner: Chief Financial Officer / Corporate Controller

Wave 4 (M 9-18): Strategic Bets & Autonomous Enterprise CoE

Deploy complex multi-system knowledge graphs and establish client's self-sustaining AI Center of Excellence.

Core Initiatives: Enterprise Knowledge Graph & Cross-Repo Developer Search Engine; Self-Sustaining AI Center of Excellence (CoE) Operationalization

Investment: ₹45 - ₹65 Lakhs (Professional Services) |
Owner: Enterprise AI Steering Committee & CTO

4.2 Workforce Enablement & Change Adoption Suite (The 3-Track Curriculum)

To ensure high employee adoption and eliminate fear of displacement, Nisol AI embeds a structured 3-Track Enablement Curriculum into the transformation engagement:

Track 1: AI Foundations
ALL-HANDS (2× 90-MIN SESSIONS)

- Demystifying generative AI & prompt mechanics
- Acceptable use policy & zero-PII data rules
- The "Future of Work": AI as a cognitive multiplier

Track 2: The New Way of Working
DEPT CHAMPIONS (4× 2-HR LABS)

- Hands-on sandbox labs with department RAG tools
- Few-shot prompt calibration & output validation
- Exception escalation & human-in-the-loop sign-off

Track 3: Leading AI Teams
LEADERSHIP & MANAGERS (HALF-DAY)


- Measuring automation ROI & team velocity
- Managing the 4 Employee Archetypes (Skeptics to Builders)
- Reallocating recovered capacity into strategic growth

4.3 Organizational Change Management (OCM) & RACI Governance Matrix

Transformation Initiative	Responsible (R)	Accountable (A)	Consulted (C)	Informed (I)
AI Strategic Roadmap & Capital Allocation	Chief AI Officer / Advisory Lead	CEO & Executive Committee	Business Unit Heads	Company-wide
AI Proxy Gateway & Data Privacy (PII) Guardrails	Security Architect / Data Ops	CISO	Legal & Compliance	All Engineering Staff
Departmental Use Case Rollout (Support/Sales/Ops)	AI Product Manager & Dev Lead	Department Head	Frontline Super-Users	Impacted Teams
Continuous Model Monitoring & Governance Audits	MLOps Engineer	AI Steering Committee	Internal Audit	Executive Leadership

5. ROI Analysis, Sensitivity & Financial Projection

5.1 5-Year Cumulative Financial Benefit vs. Investment



5.2 Balance Sheet Impact: Hard vs. Soft Savings Decomposition

Benefit Tier	Description & Operating Mechanism	Annual Value	CFO Realization Weight
Tier 1: Hard Cash Savings	Direct reduction in manual QA contractor spend, automated invoice processing, and eliminated external agency drafting fees.	₹1.60 Cr	100% (Direct EBITDA)
Tier 2: Risk & Working Capital	Avoided regulatory non-compliance fines (DPDP Act), zero customer PII exfiltration, and reduced invoice payment reconciliation errors.	₹0.95 Cr	75% (Risk-Adjusted)
Tier 3: Productivity & Soft Value	Senior engineering and managerial capacity recovered for high-velocity revenue-generating feature delivery and strategic sales growth.	₹0.65 Cr	50% (Reinvestment)

5.3 3-Scenario Sensitivity Stress Test (Base / Conservative / Optimistic)

CFO & Board stress-test evaluating program financial resilience under varying user adoption rates (75% to 125%) and cost variances at a 10% discount rate.

Scenario	Adoption Rate	Annual Savings	5-Year Net Benefit	5-Yr ROI %	Payback Period	Net Present Value (NPV)
Conservative Case	75%	₹6.60 Crore	₹24.90 Crore	431%	7.8 Mo	₹23.90 Crore
Base Case	100%	₹8.80 Crore	₹36.49 Crore	634%	5.8 Mo	₹33.20 Crore
Optimistic Case	125%	₹11.00 Crore	₹48.17 Crore	856%	4.5 Mo	₹43.16 Crore

CONFIDENTIAL

6. Solution Blueprints for High-Impact Priority Initiatives

Architectural specifications and engineering blueprints defining target components, system dependencies, and phased delivery plans for top-ranked transformation initiatives.

6.1 Blueprint: AI-Driven Automated QA Test Case Generation & Execution

Software Engineering & QA

The architecture leverages an event-driven GitHub Action triggered on Pull Requests. A Python-based parser utilizes Tree-sitter AST (Abstract Syntax Tree) to analyze code diffs, identifying modified UI components, API endpoints, and routing logic. This structural diff, along with historical test schemas, is packaged into a context-rich prompt and sent to Claude 3.5 Sonnet. The LLM generates syntactically correct Playwright/Cypress test scripts. These scripts are dynamically executed inside a sandboxed container, and execution results (including failure logs) are posted back to the PR as comments for developer review.

TECHNOLOGY STACK

Playwright, Cypress, Tree-sitter AST, Claude 3.5 Sonnet API, GitHub Actions, Docker (Sandboxed Test Environment), Node.js / TypeScript

TIMELINE & INVESTMENT

5 Weeks | ₹8 - ₹12 Lakhs | ROI: 280%

KEY OBJECTIVES

- Reduce manual regression testing cycles from 48-72 hours to under 2 hours.
- Automate the generation of Playwright and Cypress end-to-end test scripts directly from pull request code diffs.
- Integrate autonomous QA agents into the existing GitHub Actions CI/CD pipeline to catch regression bugs pre-merge.

6.2 Blueprint: AI Proposal & Technical RFP Response Generator

Sales & Business Development

A Hybrid RAG (Retrieval-Augmented Generation) pipeline built on FastAPI. Document ingestion pipelines systematically parse past winning proposals, security questionnaires (SOC2, ISO), and rate cards from SharePoint and Salesforce, chunking and indexing them using a hybrid search system (BM25 lexical search combined with dense vector embeddings in Qdrant). When a new RFP is uploaded, the system extracts the questions, performs hybrid retrieval to find the most relevant historical answers, and uses an LLM (Claude 3.5 Sonnet) to synthesize tailored, context-aware responses. A Python Docx Generator compiles the responses into a formatted corporate template.

TECHNOLOGY STACK

Hybrid RAG (BM25 + Qdrant Vector DB), FastAPI (Backend API), React (Frontend UI for Sales/Bid Managers), Salesforce CRM API, SharePoint Cloud Document Store API, LangChain / LlamaIndex, Python docx / PyPDF2

TIMELINE & INVESTMENT

6 Weeks | ₹10 - ₹15 Lakhs | ROI: 340%

KEY OBJECTIVES

- Reduce technical RFP bid preparation time from 20+ hours to under 4 hours per submission.
- Consolidate and index unstructured organizational knowledge across SharePoint, legacy CRMs, and security audits.
- Generate high-quality, 80% complete draft proposals that accurately reflect Novatech's technical capabilities, security posture, and rate cards.

6.3 Blueprint: Autonomous PMO Operational Status & Flight-Deck Reporter

Corporate Strategy & PMO

An autonomous PMO orchestration engine built in Python. The system schedules weekly cron jobs to pull data from Jira (sprint metrics, issue statuses, epic progress), Slack (team channel discussions, blocker mentions), and SharePoint (project charters, milestone schedules). A lightweight LLM processing pipeline filters, correlates, and synthesizes this raw data into structured executive summaries. The system identifies key risks (e.g., scope creep, stagnant tickets, resource overallocation) using rule-based metrics combined with semantic analysis of Slack/Jira comments. The final report is formatted into markdown, HTML, and PDF, and distributed via Slack webhooks and email.

TECHNOLOGY STACK

Jira REST API, Slack Webhooks & Slack API, Python Orchestrator (Celery / Redis for scheduling), SharePoint Cloud Document Store API, Claude 3.5 Sonnet (for semantic synthesis and summarization), WeasyPrint (HTML-to-PDF conversion)

TIMELINE & INVESTMENT

4 Weeks | ₹6 - ₹9 Lakhs | ROI: 220%

KEY OBJECTIVES

- Eliminate 6-8 hours weekly of manual status aggregation per Project Manager.
- Synthesize real-time sprint velocities, milestone blockers, and burn rates into standardized executive digests.
- Provide proactive risk alerts regarding project delays and resource bottlenecks.

CONFIDENTIAL

7. Path to Value: Recommended Implementation Roadmap & Investment

Grounded in our discovery findings, Nisol AI has structured a consultative implementation roadmap. Delivery velocity is calibrated to our audited **Single-Pod Capacity Benchmark** to guarantee high execution quality without taxing internal client teams.

7.1 Recommended AI Initiatives (Priority Order)

Rank	Initiative Name	Business Value	Complexity	Duration	Priority
1	AI-Driven Automated QA Test Case Generation	High (Score: 92)	Low	5 Weeks	Wave 1 Anchor
2	AI Technical Proposal & RFP Response Generator	High (Score: 88)	Low	6 Weeks	Wave 1 Quick Win
3	Autonomous PMO Operational Status Reporter	High (Score: 82)	Low	4 Weeks	Wave 1 Quick Win
4	Intelligent Invoice & Vendor Document OCR	High (Score: 90)	Medium	10 Weeks	Wave 2 Priority
5	Enterprise Knowledge Graph & Multi-Repo Search	High (Score: 94)	High	20 Weeks	Wave 3 Strategic

7.2 Solution Blueprint Investment Estimates

Initiative	Duration	Investment Range (INR)	Annual Value Realization
Automated QA Test Generation	5 Weeks	₹18,00,000 – ₹25,00,000	₹32 Lakhs/year
AI Technical Proposal & RFP Generator	6 Weeks	₹18,00,000 – ₹25,00,000	₹38 Lakhs/year
Autonomous PMO Status Reporter	4 Weeks	₹14,00,000 – ₹20,00,000	₹18 Lakhs/year
Intelligent Invoice OCR Processing	10 Weeks	₹28,00,000 – ₹38,00,000	₹42 Lakhs/year

Initiative	Duration	Investment Range (INR)	Annual Value Realization
Enterprise Knowledge Graph & Search	20 Weeks	₹45,00,000 – ₹65,00,000	₹65 Lakhs/year

7.3 Cost Inclusions & Commercial Terms (Training Included, Infra Pass-Through)

✓ **What Is INCLUDED in Nisol Professional Fees**

- **Dedicated Delivery Pod:** AI Lead, Full-Stack Engineers & Data Specialists (1 Pod = 1 Initiative per 12–14 wk cycle).
- **Custom Solution Engineering:** RAG pipelines, API orchestration, microservices, and evaluation harnesses.
- **Workforce Enablement & Employee Training Program:**
 - *AI Readiness & Demystification:* 2× 90-min All-Hands sessions so employees do not start from zero.
 - *The New Way of Working:* 4× 2-hr hands-on sandbox labs for department super-users.
 - *Leading AI Teams:* 1x half-day seminar for managers on workload re-allocation.
- **Continuous Quality Gates:** Production readiness audits, red-teaming, and post-rollout hypercare.

⚠ **What Is EXCLUDED (Direct Client Pass-Through)**

- **Cloud Infrastructure Compute:** All AWS, Microsoft Azure, or GCP VPC virtual machines, serverless containers, and storage.
- **Foundation Model Token APIs:** LLM inference tokens (OpenAI GPT-4o, Anthropic Claude 3.5, Google Gemini 1.5).
- **Vector Database Hosting:** Managed vector DB subscriptions (Pinecone, Qdrant Cloud, Weaviate).
- **Third-Party Enterprise Licenses:** API connectors and seats for CRM/ERP systems (Salesforce, SAP, Zendesk, Jira).
- *Scope Note:* All environment instances across Development, Staging, UAT, and Production are provisioned inside Novatech Systems Ltd.'s tenant and billed directly to Novatech Systems Ltd..

Delivery Capacity Benchmark: The Single-Pod Principle

To guarantee rigorous engineering quality and protect internal engineering capacity, **one dedicated Nisol delivery pod executes exactly one strategic initiative per 12-to-14 week cycle**. Concurrent initiative execution requires authorization of independent parallel pods.

8. Terms & Conditions & Executive Acceptance

This document constitutes the official Enterprise AI Transformation Assessment and Phased Implementation Strategy formulated by Nisol AI Advisory Services for **Novatech Systems Ltd..** Signature below signifies executive acceptance of the assessment findings and authorization to initiate Phase 1 engagement scopes under the agreed Single-Pod delivery benchmark and pass-through commercial terms.

For: Novatech Systems Ltd.

Authorized Executive Signature & Date

For: Nisol AI Advisory

Managing Partner / Lead AI Advisor Signature & Date

CONFIDENTIAL

BOARD INVESTMENT MEMORANDUM

CAPITAL ALLOCATION & BUSINESS CASE FOR ENTERPRISE AI TRANSFORMATION

Prepared for the Board of Directors & Chief Financial Officer

CONFIDENTIAL

ORGANIZATION

Novatech Systems Ltd.

INDUSTRY SECTOR

Technology, SaaS & IT Services

DECISION GATE

Phase 1 Capital Authorization

DATE OF ISSUE

August 31, 2026

1. Executive Investment Teaser & Key Returns


This memorandum requests board authorization for a phased capital allocation of **₹4.26 Crore** to execute the Enterprise AI Transformation Program at **Novatech Systems Ltd.**. The program targets operational throughput acceleration and manual workflow automation across Customer Support, Engineering QA, and Financial reconciliation.

CAPITAL INVESTMENT	ANNUAL SAVINGS	5-YR NET BENEFIT	PAYBACK PERIOD
₹4.26 Crore	₹8.80 Crore	₹36.49 Crore	5.8 Months
Phase 1-2 Total	Recurring run-rate	Net of all costs	IRR: 44.5%

Balance Sheet Impact: Hard vs. Soft Savings Decomposition			
Benefit Tier	Description	Annual Value	CFO Realization Weight
Tier 1: Hard Cash Savings	Direct reductions in manual QA contractor spend & invoice processing labor	₹1.60 Cr	100% (Direct EBITDA)
Tier 2: Risk & Working Capital	Avoided regulatory fines, zero PII exfiltration, and reduced invoice payment errors	₹0.95 Cr	75% (Risk-Adjusted)
Tier 3: Productivity & Soft Value	Senior developer capacity recovered for revenue-generating feature delivery	₹0.65 Cr	50% (Reinvestment)

Strategic Imperative: Cost of Inaction & Valuation Multiple Premium	
In the Technology, SaaS & IT Services sector, top quartile competitors are investing 4.5% of revenue into automation. Delaying enterprise AI adoption introduces an estimated annual opportunity deficit of ~₹8.80 Crore in manual labor drag. Furthermore, empirical McKinsey and Morgan Stanley data demonstrates that enterprises with scaled AI execution trade at an average 15–25% valuation multiple premium over sector peers due to superior operating leverage.	

2. 5-Year Cash Flow & Sensitivity Stress Test



3-Scenario Sensitivity Stress Test (10% Discount Rate)

Stress-test demonstrating program financial viability even if employee adoption drops to 75% and infrastructure maintenance expands by +15%.

Scenario	Adoption Rate	Annual Savings	5-Year Net Benefit	5-Yr ROI %	Payback Period	Net Present Value (NPV)
Conservative Case	75%	₹6.60 Crore	₹24.90 Crore	431%	7.8 Mo	₹23.90 Crore
Base Case	100%	₹8.80 Crore	₹36.49 Crore	634%	5.8 Mo	₹33.20 Crore
Optimistic Case	125%	₹11.00 Crore	₹48.17 Crore	856%	4.5 Mo	₹43.16 Crore

3. Formal Board Resolution & Commercial Terms

Commercial Term: Direct Client Infrastructure Pass-Through & Training Inclusions

- **Infrastructure Pass-Through:** The investment authorization requested herein covers Nisol AI Professional Advisory and Engineering Services exclusively. All third-party cloud compute (AWS/Azure/GCP VPCs), foundation model token API consumption, vector database hosting, and enterprise SaaS connectors across Development, Staging, UAT, and Live environments shall be provisioned directly within Novatech Systems Ltd.'s enterprise accounts and paid directly by Novatech Systems Ltd..
- **Workforce Enablement Included:** The authorized professional fees include the full 3-Track Workforce Enablement Curriculum (Track 1: All-Hands AI Foundations, Track 2: Department Champions Sandbox Labs, Track 3: Management Seminar) to ensure staff do not start from zero and achieve rapid operational adoption.

RESOLVED THAT, the Board of Directors of **Novatech Systems Ltd.** hereby approves the allocation of **₹4.26 Crore** for Phase 1 of the Enterprise AI Transformation Program under the Single-Pod governance model, authorizing executive leadership to execute the initial scope of work with Nisol AI Advisory.

For the Board of Directors: Novatech Systems Ltd.

Chairman / Chief Financial Officer (CFO)
Signature & Date

For: Nisol AI Advisory

Managing Partner / Lead AI Advisor Signature & Date

TECHNICAL ARCHITECTURE DOSSIER

DATA STRATEGY & HYBRID VECTOR LAKEHOUSE BLUEPRINT

Prepared for the Chief Technology Officer & Data Engineering Leadership

CONFIDENTIAL

ORGANIZATION

Novatech Systems Ltd.

OVERALL DATA READINESS

62 / 100

ESTIMATED DATA PREP BUDGET

₹22 - ₹32 Lakhs

DATE OF ISSUE

August 31, 2026

1. Enterprise Data Quality Assessment (5 Dimensions)

McKinsey research demonstrates that 87% of AI initiative failures stem from unresolved data pipeline and hygiene bottlenecks. This evaluation establishes the baseline data quality posture for **Novatech Systems Ltd..**

Dimension	Quality Score	Status	Audit Findings & Gap Analysis
Completeness	57%	Needs Attention	Legacy transactional records have variable mandatory field population; document archives contain unindexed scanned PDFs.
Accuracy	61%	Needs Attention	High core financial and operational precision, but cross-system customer identity records require deduplication.
Timeliness	53%	Healthy	Batch ETL sync intervals (daily/weekly) create lag; real-time event streaming is required for autonomous agent response.
Consistency	51%	Needs Attention	Discrepant schema naming conventions across sales CRM vs ERP databases require canonical data modeling.
Accessibility	55%	Critical Blocker	Data locked in siloed department repositories without unified semantic vector embeddings or secure REST/GraphQL endpoints.

Data Preparation Budget Allocation

Data cleansing, OCR ingestion, vector embedding extraction, and metadata tagging represent **35% (₹22 - ₹32 Lakhs)** of the Phase 1 transformation budget. Allocating these funds early prevents expensive hallucination remediation during agent rollout.

2. Target Architecture Reference Pattern (Hypothesis)

Based on initial discovery findings across disparate document repositories and ERP ledgers, Nisol AI presents the following **Target Architecture Hypothesis**. *Note: Specific vector topologies (pgvector vs. Qdrant vs. native Snowflake/Databricks vectors) and CDC streaming mechanisms are subject to formal validation during the Phase 1 Technical Discovery Spike based on Novatech Systems Ltd.'s cloud VPC security boundaries and data residency mandates.*

Commercial Term: Client Infrastructure Pass-Through

All vector database hosting, cloud compute instances (AWS/Azure/GCP), and foundation model token consumption are direct pass-through expenses provisioned within Novatech Systems Ltd.'s enterprise cloud tenant. Nisol AI fees reflect custom pipeline engineering, prompt orchestration, and evaluation harnesses exclusively.

Vector Indexing & Embedding Model

Dense vector embeddings (BGE-M3 / OpenAI text-embedding-3-large) combined with sparse BM25 keyword indices inside pgvector / Qdrant with HNSW cosine similarity search.

Change Data Capture (CDC) Pipeline

Real-time Debezium / Kafka streaming connectors replicating modified CRM, ERP, and ticket tables into vector stores with automated chunk invalidation.

Recommended Data Engineering Roadmap

- Build a centralized Hybrid Vector Lakehouse (e.g. pgvector / Qdrant with hybrid BM25 search)
- Deploy an automated Document Ingestion & OCR pipeline with chunk-level metadata tagging
- Implement real-time Change Data Capture (CDC) pipelines from ERP/CRM into analytical vector stores
- Configure granular Role-Based Access Control (RBAC) filtering at the embedding chunk level

PILOT DECISION GATE & ACCEPTANCE PROTOCOL

PoC EVALUATION FRAMEWORK
& SCALABILITY DOSSIER

Pre-Execution Governance Protocol: Hypotheses, SLA Benchmarks & Production Gate Criteria

CONFIDENTIAL

ORGANIZATION

Novatech Systems Ltd.

STANDARD PILOT SPRINT

4-to-6 Week Controlled
Sandbox

GOVERNANCE GATE

GATE 0: PILOT SCOPE
AUTHORIZATION

DATE OF ISSUE

August 31, 2026

1. Pilot Hypotheses, Acceptance Benchmarks & Evaluation Harness

Prior to authorizing live production deployment for any AI capability, Nisol AI mandates a controlled **4-to-6 week sandbox pilot**. This document establishes the prospective quantitative criteria, synthetic stress-test harness, and institutional decision gates required for Phase 1 sign-off.

TARGET CITATION PRECISION

> 95.0%

Zero tolerance for uncited hallucinations

TARGET INFERENCE LATENCY

< 2.0s

Interactive UI SLA at p95

TARGET USER TASK REDUCTION

> 40.0%

Measured vs. human baseline

Candidate Pilot Workstreams & Pre-Agreed Success Thresholds

Candidate Pilot Workstream	Hypothesis to Validate	Required Acceptance Threshold	Evaluation Harness
Automated QA Test Generation Engineering & DevOps	Agent auto-generates executable Playwright tests from PR diffs	≥ 85% valid syntactical execution; zero broken existing tests	Golden CI/CD Suite
Intelligent Invoice OCR Parser Finance & Accounting	Vision LLM extracts 15 line items and matches ERP PO numbers	≥ 95% field extraction accuracy across 100 historical invoices	ERP Ledger Match
Enterprise SOP Knowledge Hub HR & Operations	RAG engine answers policy queries with exact PDF handbook citations	≥ 98% factual precision across 200 synthetic challenge queries	RAG Triad Benchmark

Institutional Safeguard: Single-Pod / Single-Initiative Governance

Only **one primary pilot workstream** shall be deployed into active testing per 12-week pod cycle. Parallel pilot testing requires the authorization of independent dedicated squads to protect internal engineering review capacity.

2. Production Rollout Gate Protocol & Authorization Sign-off

At the conclusion of the 4-week sandbox testing window, the joint Steering Committee reviews empirical metrics against the four mandatory decision gates:

Decision Gate	Evaluation Method	Pass Condition	Accountable Lead
Gate 1: Accuracy & Fidelity	Automated evaluation against 200 ground-truth golden test cases	$\geq 95.0\%$ Accuracy	AI Solutions Architect
Gate 2: Security & PII Redaction	Adversarial red-teaming and prompt-injection stress testing	Zero PII Leaks	CISO / InfoSec Sponsor
Gate 3: End-User Usability	Department super-user evaluation across 5 business workflows	$\geq 80\%$ User Satisfaction	Department Business Sponsor
Gate 4: Unit Economics	Inference token cost per completed workflow calculation	Within Target OPEX Budget	Corporate Controller

Formal Authorization Protocol: Gate 0 Approval

By executing below, Novatech Systems Ltd. authorizes Nisol AI to initiate the Sandbox Pilot Sprint under the terms and quantitative acceptance criteria outlined in this dossier.

For: Novatech Systems Ltd. Evaluation Sponsor

Product Leader / VP Sponsor Signature & Date

For: Nisol AI Advisory

Lead AI Solutions Architect Signature & Date